

AI / ML Engineering

Delivering measurable client value

Ruibin Chen | Technical Lead & AI Engineer | ~12 years | production AI/ML, outcome-focused



Energy Agent

Agentic AI customer service

-50%

human-agent tickets



Australia Post CV

Warehouse safety detection

Fewer

accidents & unsafe acts



Adora AI

Energy use, by appliance

New

revenue: tenants want to buy



TallyPlus-AI

Internal AI knowledge tool

Faster

delivery across the org

Energy Agent — automating customer service



The problem

A major energy retailer's support team carried high volumes of routine billing, plan and complaint tickets — slow to clear and costly to staff.



What we delivered

An AI agent that resolves routine tickets end-to-end (billing, plans, complaints, bookings), with a voice channel and automatic call summaries. Serves ~100K active users of a ~3M base.



The return

Customer-raised human-agent tickets fell ~50% as the agent absorbed them; average ticket resolution dropped from ~48h to ~36h; customer satisfaction edged up.

-50%

human-agent tickets

48h → 36h

ticket resolution time

Scale & my role

~100K active users
of a ~3M customer base

Led design and delivery end-to-end; built on AWS (Bedrock, agentic orchestration).

"The LLM is one small piece — the engineering is everything around it."

Range

Client value across domains



Australia Post — computer vision

Safety & loss reduction

Fewer accidents

Real-time detection of unsafe behaviour in the warehouse. After release, accidents and unsafe acts dropped — safer staff and lower loss for the business.

96% detection accuracy. YOLOv3, deployed on AWS SageMaker.



Adora AI

Customer experience → new revenue

New sales interest

Breaks a household's energy use down by appliance, lifting customer experience. Demoed to other SaaS tenants — strong buying interest in the feature.

Serves 3M customers. NILM models on GCP.



TallyPlus-AI

Org-wide productivity

Faster delivery

Internal AI knowledge tool over 6,800+ docs. Lifted delivery velocity — and adoption spread beyond devs to PM and QA, all using AI to speed up.

6,800+ docs / ~23,000 chunks. RAG on AWS Bedrock.